



# **Detailed Project Report(DPR)**

**SUBMITTED BY**

**RASHID ALI**

(IMAGE CAPTIONING PROJECT)

---

INEURON INTERNSHIP

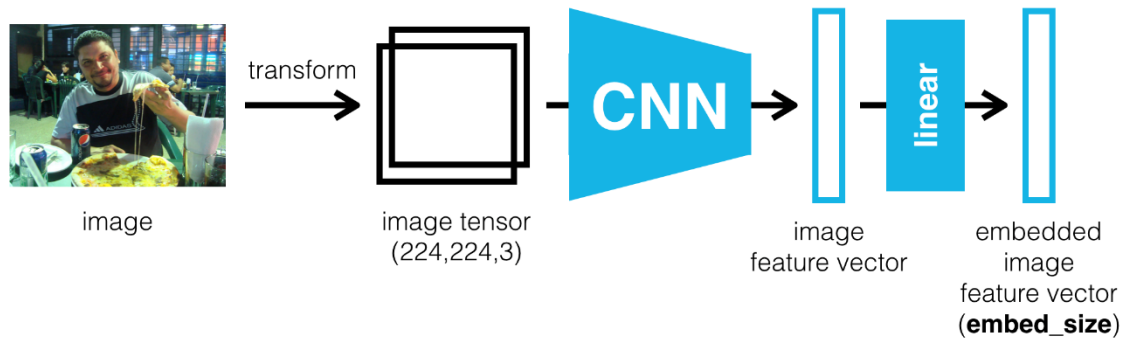
---

Domain : Security and Safety, Surveillance

**IMAGE - CAPTIONING**

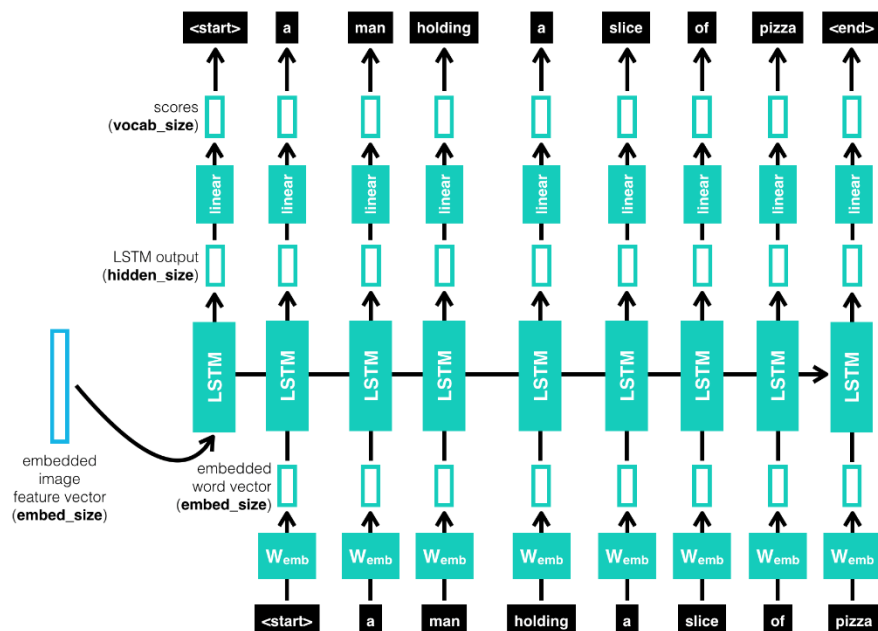
# Explaining Encoder and Decoder

## Encoder:



For Encoder we use ResNet-50. ResNet-50 is a convolutional neural network that is trained on more than a million images from the ImageNet database. The network is 50 layers deep and can classify images into 1000 object categories, such as keyboard, mouse, pencil, and many animals. As a result, the network has learned rich feature representations for a wide range of images. The network has an image input size of 224-by-224.

## Decode:



For Decoder we use LSTM. Long Short-Term Memory (LSTM) networks are a modified version of recurrent neural networks, which makes it easier to remember past data in memory. The vanishing gradient problem of RNN is resolved here. LSTM is well-suited to classify, process and predict time series given time lags of unknown duration. It trains the model by using back-propagation

# Table of Contents

- Let's do it...
  - Step 1 – Importing required libraries for Image Captioning.
  - Step 2 – Extract features from the images using VGG-16.
  - Step 3 – Load, Clean and Save image descriptions.
  - Step 4 – Load train and test image features and descriptions.
  - Step 5 – Getting descriptions in shape.
  - Step 6 – Functions to generate data and create model.
  - Step 7 – Creating and training Image Captioning model.
  - Step 8 – Prediction and Evaluation functions for Image Captioning model.
  - Step 9 – Evaluating Image Captioning model.
  - Step 10 – Live Image Captioning.

# Image Caption Generator using Python | Flickr 8k Dataset

Image caption generator is a process of recognizing the context of an image and annotating it with relevant captions using deep learning, and computer vision. This is an advanced deep learning project where more than one model must be used for analysis and preprocessing the data to obtain the results.

In this project tutorial, we will build an image caption generator to load a random image and give some captions describing the image. We will use Convolutional Neural Network (CNN) for image feature extraction and Long Short-Term Memory Network (LSTM) for Natural Language Processing (NLP).

## Dataset Information

The objective of the project is to predict the captions for the input image. The dataset consists of 8k images and 5 captions for each image. The features are extracted from both the image and the text captions for input.

The features will be concatenated to predict the next word of the caption. CNN is used for image and LSTM is used for text. BLEU Score is used as a metric to evaluate the performance of the trained model.

## Import Modules

First, we have to import all the basic modules we will be needing for this project

```
import os
import pickle
import numpy as np
from tqdm.notebook import tqdm
from tensorflow.keras.applications.vgg16 import VGG16, preprocess_input
from tensorflow.keras.preprocessing.image import load_img, img_to_array
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Model
from tensorflow.keras.utils import to_categorical, plot_model
from tensorflow.keras.layers import Input, Dense, LSTM, Embedding, Dropout, add
```

- **os** - used to handle files using system commands.
- **pickle** - used to store numpy features extracted

- **numpy** - used to perform a wide variety of mathematical operations on arrays
- **tqdm** - progress bar decorator for iterators. Includes a default range iterator printing to stderr.
- **VGG16, preprocess\_input** - imported modules for feature extraction from the image data
- **load\_img, img\_to\_array** - used for loading the image and converting the image to a numpy array
- **Tokenizer** - used for loading the text as convert them into a token
- **pad\_sequences** - used for equal distribution of words in sentences filling the remaining spaces with zeros
- **plot\_model** - used to visualize the architecture of the model through different images

## Now we must set the directories to use the data

```
BASE_DIR = '/kaggle/input/flickr8k'
WORKING_DIR = '/kaggle/working'
```

# Extract Image Features

## We have to load and restructure the model

```
# load vgg16 model
model = VGG16()
# restructure the model
model = Model(inputs=model.inputs, outputs=model.layers[-2].output)
# summarize
print(model.summary())
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16\_weights\_tf\_dim\_ordering\_tf\_kernels.h5 553467904/553467096  
[=====] - 3s 0us/step 553476096/553467096  
[=====] - 3s 0us/step Model: "model"

Output Shape	Param #	Layer (type)
(InputLayer) [(None, 224, 224, 3)]	0	input_1
block1_conv1 (Conv2D) (None, 224, 224, 64)	1792	
block1_conv2 (Conv2D) (None, 224, 224, 64)	36928	
(MaxPooling2D) (None, 112, 112, 64)	0	block1_pool
block2_conv1 (Conv2D) (None, 112, 112, 128)	73856	
block2_conv2 (Conv2D) (None, 112, 112, 128)	147584	
(MaxPooling2D) (None, 56, 56, 128)	0	block2_pool
block3_conv1 (Conv2D) (None, 56, 56, 256)	295168	

block3_conv2 (Conv2D) (None, 56, 56, 256) 590080	
block3_conv3 (Conv2D) (None, 56, 56, 256) 590080	
(MaxPooling2D) (None, 28, 28, 256) 0	block3_pool
block4_conv1 (Conv2D) (None, 28, 28, 512) 1180160	
block4_conv2 (Conv2D) (None, 28, 28, 512) 2359808	
block4_conv3 (Conv2D) (None, 28, 28, 512) 2359808	
(MaxPooling2D) (None, 14, 14, 512) 0	block4_pool
block5_conv1 (Conv2D) (None, 14, 14, 512) 2359808	
block5_conv2 (Conv2D) (None, 14, 14, 512) 2359808	
block5_conv3 (Conv2D) (None, 14, 14, 512) 2359808	
(MaxPooling2D) (None, 7, 7, 512) 0	block5_pool
(Flatten) (None, 25088) 0	flatten
(None, 4096) 102764544	fc1 (Dense)
(None, 4096) 16781312	fc2 (Dense)
===== Total	
params: 134,260,544 Trainable params: 134,260,544 Non-trainable params: 0	
	None

- Fully connected layer of the VGG16 model is not needed, just the previous layers to extract feature results.
- By preference you may include more layers, but for quicker results avoid adding the unnecessary layers.

## Now we extract the image features and load the data for preprocess

*# extract features from image*

**features = {}**

directory = os.path.join(BASE\_DIR, 'Images')

```
for img_name in tqdm(os.listdir(directory)):
    # load the image from file
    img_path = directory + '/' + img_name
    image = load_img(img_path, target_size=(224, 224))
    # convert image pixels to numpy array
    image = img_to_array(image)
    # reshape data for model
    image = image.reshape((1, image.shape[0], image.shape[1], image.shape[2]))
    # preprocess image for vgg
    image = preprocess_input(image)
    # extract features
```

```

feature = model.predict(image, verbose=0)
# get image ID
image_id = img_name.split('.')[0]
# store feature
features[image_id] = feature
    • Dictionary 'features' is created and will be loaded with the extracted features of image data
    • load_img(img_path, target_size=(224, 224)) - custom dimension to resize the image when loaded to the array
    • image.reshape((1, image.shape[0], image.shape[1], image.shape[2])) - reshaping the image data to preprocess in a RGB type image.
    • model.predict(image, verbose=0) - extraction of features from the image
    • img_name.split('.')[0] - split of the image name from the extension to load only the image name.
# store features in pickle
pickle.dump(features, open(os.path.join(WORKING_DIR, 'features.pkl'), 'wb'))
    • Extracted features are not stored in the disk, so re-extraction of features can extend running time
    • Dumps and store your dictionary in a pickle for reloading it to save time

# load features from pickle
with open(os.path.join(WORKING_DIR, 'features.pkl'), 'rb') as f:
    features = pickle.load(f)
    • Load all your stored feature data to your project for quicker runtime

```

# Load the Captions Data

## Let us store the captions data from the text file

```

with open(os.path.join(BASE_DIR, 'captions.txt'), 'r') as f:
    next(f)
    captions_doc = f.read()

```

## Now we split and append the captions data with the image

```

# create mapping of image to captions
mapping = {}
# process lines
for line in tqdm(captions_doc.split('\n')):
    # split the line by comma(,)
    tokens = line.split(',')
    if len(line) < 2:
        continue
    image_id, caption = tokens[0], tokens[1:]

```

```

# remove extension from image ID
image_id = image_id.split('.')[0]
# convert caption list to string
caption = " ".join(caption)
# create list if needed
if image_id not in mapping:
    mapping[image_id] = []
# store the caption
mapping[image_id].append(caption)

```

- Dictionary 'mapping' is created with key as image\_id and values as the corresponding caption text
- Same image may have multiple captions, **if image\_id not in mapping:**  
**mapping[image\_id] = []** creates a list for appending captions to the corresponding image

## Now let us see the no. of images loaded

```

len(mapping)
8091

```

Preprocess Text Data

```

def clean(mapping):
    for key, captions in mapping.items():
        for i in range(len(captions)):
            # take one caption at a time
            caption = captions[i]
            # preprocessing steps
            # convert to lowercase
            caption = caption.lower()
            # delete digits, special chars, etc.,
            caption = caption.replace('[^A-Za-z]', '')
            # delete additional spaces
            caption = caption.replace('\s+', ' ')
            # add start and end tags to the caption
            caption = 'startseq ' + " ".join([word for word in caption.split() if
len(word)>1]) + ' endseq'
            captions[i] = caption
            • Defined to clean and convert the text for quicker process and better results

```

## Let us visualize the text before and after cleaning

```

# before preprocess of text
mapping['1000268201_693b08cb0e']

```



['A child in a pink dress is climbing up a set of stairs in an entry way .', 'A girl going into a wooden building .', 'A little girl climbing into a wooden playhouse .', 'A little girl climbing the stairs to her playhouse .', 'A little girl in a pink dress going into a wooden cabin .']

*# preprocess the text*

```
clean(mapping)
```

*# after preprocess of text*

```
mapping['1000268201_693b08cb0e']
```

['startseq child in pink dress is climbing up set of stairs in an entry way endseq', 'startseq girl going into wooden building endseq', 'startseq little girl climbing into wooden playhouse endseq', 'startseq little girl climbing the stairs to her playhouse endseq', 'startseq little girl in pink dress going into wooden cabin endseq']

- Words with one letter was deleted
- All special characters were deleted
- 'startseq' and 'endseq' tags were added to indicate the start and end of a caption for easier processing

## Next we will store the preprocessed captions into a list

```
all_captions = []
```

```
for key in mapping:
```

```
    for caption in mapping[key]:  
        all_captions.append(caption)
```

```
len(all_captions)
```

```
40455
```

- No. of unique captions stored

## Let us see the first ten captions

```
all_captions[:10]
```

['startseq child in pink dress is climbing up set of stairs in an entry way endseq', 'startseq girl going into wooden building endseq', 'startseq little girl climbing into wooden playhouse endseq', 'startseq little girl climbing the stairs to her playhouse endseq', 'startseq little girl in pink dress going into wooden cabin endseq', 'startseq black dog and spotted dog are fighting endseq', 'startseq black dog and tri-colored dog playing with each other on the road endseq', 'startseq black dog and white dog with brown spots are staring at each other in the street endseq', 'startseq two dogs of different breeds looking at each other on the road endseq', 'startseq two dogs on pavement moving toward each other endseq']

## Now we start processing the text data

*# tokenize the text*

```
tokenizer = Tokenizer()
```

```
tokenizer.fit_on_texts(all_captions)
```

```
vocab_size = len(tokenizer.word_index) + 1
```

```
vocab_size
8485
```

- No. of unique words

```
# get maximum length of the caption available
```

```
max_length = max(len(caption.split()) for caption in all_captions)
```

```
max_length
35
```

- Finding the maximum length of the captions, used for reference for the padding sequence.

Train Test Split

## After preprocessing the data now we will train, test and split

```
image_ids = list(mapping.keys())
split = int(len(image_ids) * 0.90)
train = image_ids[:split]
test = image_ids[split:]
```

## Explanatory example of the sequence split into pairs

```
# startseq girl going into wooden building endseq
#      X                      y
# startseq                      girl
# startseq girl                 going
# startseq girl going           into
# .....
# startseq girl going into wooden building      endseq
```

Now we will define a batch and include the padding sequence

```
# create data generator to get data in batch (avoids session crash)
```

```
def data_generator(data_keys, mapping, features, tokenizer, max_length, vocab_size,
batch_size):
```

```
    # loop over images
```

```
    X1, X2, y = list(), list(), list()
```

```
    n = 0
```

```
    while 1:
```

```
        for key in data_keys:
```

```
            n += 1
```

```
            captions = mapping[key]
```

```
            # process each caption
```

```
            for caption in captions:
```

```
                # encode the sequence
```

```

seq = tokenizer.texts_to_sequences([caption])[0]
# split the sequence into X, y pairs
for i in range(1, len(seq)):
    # split into input and output pairs
    in_seq, out_seq = seq[:i], seq[i]
    # pad input sequence
    in_seq = pad_sequences([in_seq], maxlen=max_length)
    [0]
    # encode output sequence
    out_seq = to_categorical([out_seq],
        num_classes=vocab_size)[0]
    # store the sequences
    X1.append(features[key][0])
    X2.append(in_seq)
    y.append(out_seq)
if n == batch_size:
    X1, X2, y = np.array(X1), np.array(X2), np.array(y)
    yield [X1, X2], y
    X1, X2, y = list(), list(), list()
    n = 0

```

- Padding sequence normalizes the size of all captions to the max size filling them with zeros for better results.

## Model Creation

```

# encoder model
# image feature layers
inputs1 = Input(shape=(4096,))
fe1 = Dropout(0.4)(inputs1)
fe2 = Dense(256, activation='relu')(fe1)
# sequence feature layers
inputs2 = Input(shape=(max_length,))
se1 = Embedding(vocab_size, 256, mask_zero=True)(inputs2)
se2 = Dropout(0.4)(se1)
se3 = LSTM(256)(se2)

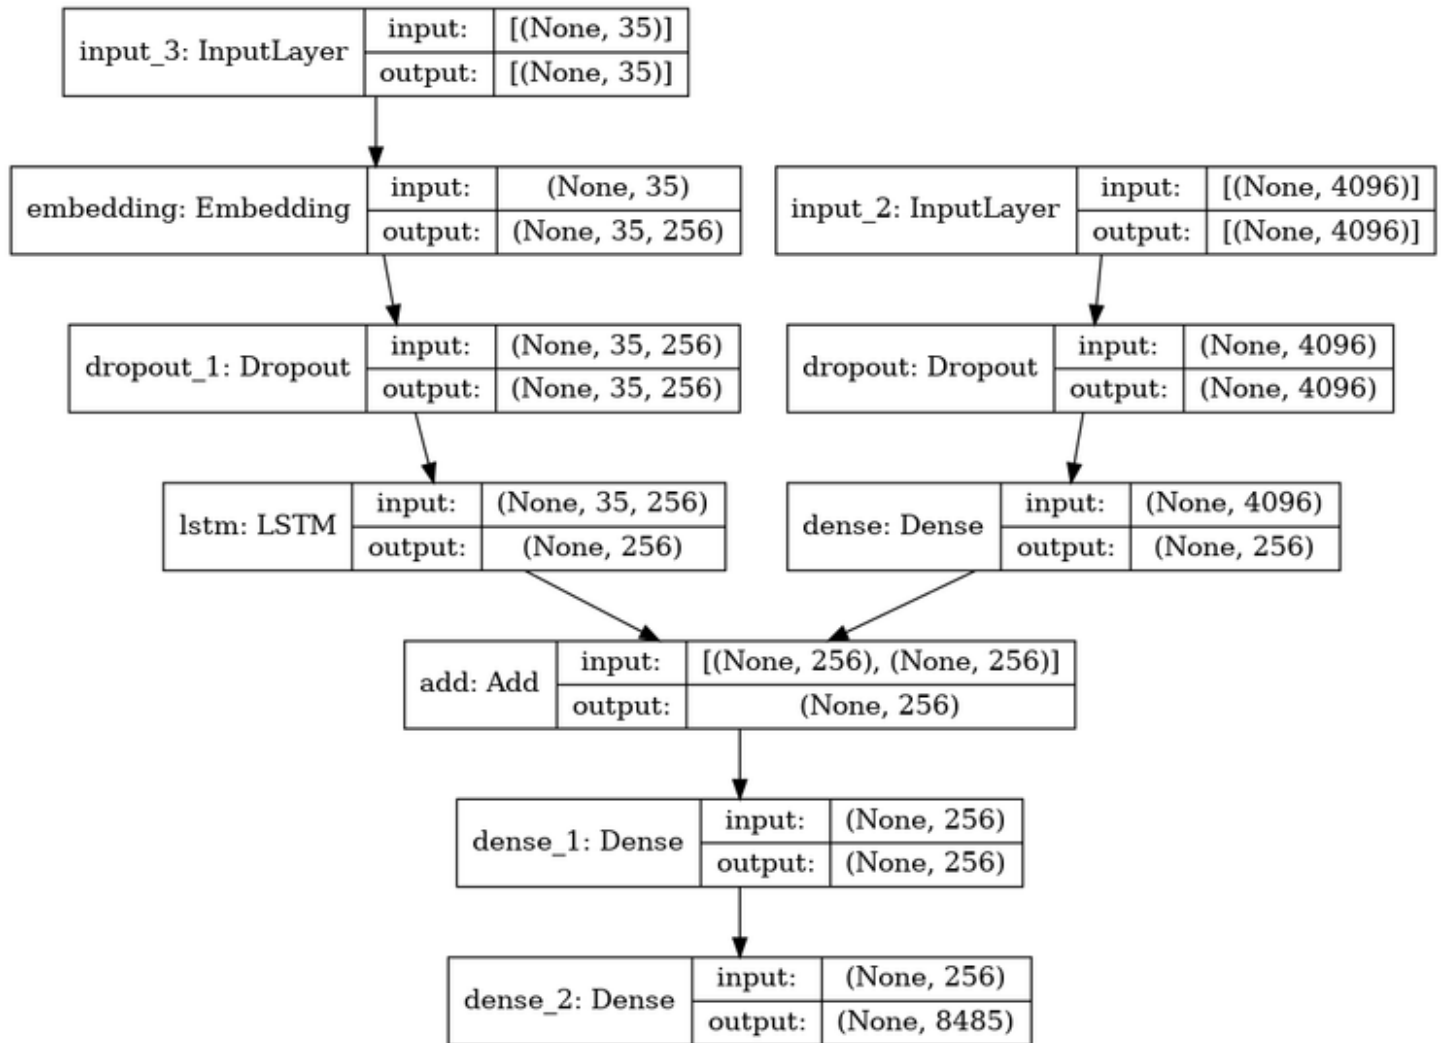
# decoder model
decoder1 = add([fe2, se3])
decoder2 = Dense(256, activation='relu')(decoder1)
outputs = Dense(vocab_size, activation='softmax')(decoder2)

model = Model(inputs=[inputs1, inputs2], outputs=outputs)
model.compile(loss='categorical_crossentropy', optimizer='adam')

# plot the model
plot_model(model, show_shapes=True)

# plot the model
plot_model(model, show_shapes=True)

```



- **shape=(4096,)** - output length of the features from the VGG model
- **Dense** - single dimension linear layer array
- **Dropout()** - used to add regularization to the data, avoiding over fitting & dropping out a fraction of the data from the layers
- **model.compile()** - compilation of the model
- **loss='sparse\_categorical\_crossentropy'** - loss function for category outputs
- **optimizer='adam'** - automatically adjust the learning rate for the model over the no. of epochs
- Model plot shows the concatenation of the inputs and outputs into a single layer
- Feature extraction of image was already done using VGG, no CNN model was needed in this step.

## Now let us train the model

*# train the model*

epochs = 20

batch\_size = 32

steps = len(train) // batch\_size

**for** i **in** range(epochs):

*# create data generator*

generator = data\_generator(train, mapping, features, tokenizer, max\_length, vocab\_size, batch\_size)

*# fit for one epoch*

model.fit(generator, epochs=1, steps\_per\_epoch=steps, verbose=1)

227/227 [=====] - 68s 285ms/step - loss: 5.2210 227/227

[=====] - 66s 291ms/step - loss: 4.0199 227/227

```
[=====] - 66s 292ms/step - loss: 3.5781 227/227
[=====] - 65s 287ms/step - loss: 3.3090 227/227
[=====] - 66s 292ms/step - loss: 3.1080 227/227
[=====] - 65s 286ms/step - loss: 2.9619 227/227
[=====] - 63s 276ms/step - loss: 2.8491 227/227
[=====] - 64s 282ms/step - loss: 2.7516 227/227
[=====] - 64s 282ms/step - loss: 2.6670 227/227
[=====] - 65s 286ms/step - loss: 2.5966 227/227
[=====] - 66s 290ms/step - loss: 2.5327 227/227
[=====] - 61s 270ms/step - loss: 2.4774 227/227
[=====] - 65s 288ms/step - loss: 2.4307 227/227
[=====] - 66s 289ms/step - loss: 2.3873 227/227
[=====] - 62s 274ms/step - loss: 2.3451 227/227
[=====] - 65s 285ms/step - loss: 2.3081 227/227
[=====] - 65s 288ms/step - loss: 2.2678 227/227
[=====] - 66s 292ms/step - loss: 2.2323 227/227
[=====] - 65s 285ms/step - loss: 2.1992 227/227
[=====] - 66s 291ms/step - loss: 2.1702
```

- **steps = len(train) // batch\_size** - back propagation and fetch the next data
- Loss decreases gradually over the iterations
- Increase the no. of epochs for better results
- Assign the no. of epochs and batch size accordingly for quicker results

## You can save the model in the working directory for reuse

```
# save the model
model.save(WORKING_DIR+'best_model.h5')
```

### Generate Captions for the Image

```
def idx_to_word(integer, tokenizer):
    for word, index in tokenizer.word_index.items():
        if index == integer:
            return word
    return None
```

- Convert the predicted index from the model into a word

```
# generate caption for an image
def predict_caption(model, image, tokenizer, max_length):
    # add start tag for generation process
    in_text = 'startseq'
    # iterate over the max length of sequence
    for i in range(max_length):
        # encode input sequence
        sequence = tokenizer.texts_to_sequences([in_text])[0]
        # pad the sequence
        sequence = pad_sequences([sequence], max_length)
        # predict next word
        yhat = model.predict([image, sequence], verbose=0)
```

```

    # get index with high probability
    yhat = np.argmax(yhat)
    # convert index to word
    word = idx_to_word(yhat, tokenizer)
    # stop if word not found
    if word is None:
        break
    # append word as input for generating next word
    in_text += " " + word
    # stop if we reach end tag
    if word == 'endseq':
        break
return in_text

```

- Caption generator appending all the words for an image
- The caption starts with 'startseq' and the model continues to predict the caption until the 'endseq' appeared

## Now we validate the data using BLEU Score

```

from nltk.translate.bleu_score import corpus_bleu
# validate with test data
actual, predicted = list(), list()

for key in tqdm(test):
    # get actual caption
    captions = mapping[key]
    # predict the caption for image
    y_pred = predict_caption(model, features[key], tokenizer, max_length)
    # split into words
    actual_captions = [caption.split() for caption in captions]
    y_pred = y_pred.split()
    # append to the list
    actual.append(actual_captions)
    predicted.append(y_pred)
    # calculate BLEU score
print("BLEU-1: %f" % corpus_bleu(actual, predicted, weights=(1.0, 0, 0, 0)))
print("BLEU-2: %f" % corpus_bleu(actual, predicted, weights=(0.5, 0.5, 0, 0)))
BLEU-1: 0.516880 BLEU-2: 0.293009

```

- BLEU Score is used to evaluate the predicted text against a reference text, in a list of tokens.
- The reference text contains all the words appended from the captions data (actual\_captions)
- A BLEU Score more than 0.4 is considered a good result, for a better score increase the no. of epochs accordingly.

# Visualize the Results

```
from PIL import Image
import matplotlib.pyplot as plt
def generate_caption(image_name):
    # load the image
    # image_name = "1001773457_577c3a7d70.jpg"
    image_id = image_name.split('.')[0]
    img_path = os.path.join(BASE_DIR, "Images", image_name)
    image = Image.open(img_path)
    captions = mapping[image_id]
    print('-----Actual-----')
    for caption in captions:
        print(caption)
    # predict the caption
    y_pred = predict_caption(model, features[image_id], tokenizer, max_length)
    print('-----Predicted-----')
    print(y_pred)
    plt.imshow(image)
```

- Image caption generator defined
- First prints the actual captions of the image then prints a predicted caption of the image

```
generate_caption("1001773457_577c3a7d70.jpg")
```

-----Actual-----

startseq black dog and spotted dog are fighting endseq

startseq black dog and tri-colored dog playing with each other on the road endseq  
startseq black dog and white dog with brown spots are staring at each other in the street  
endseq

startseq two dogs of different breeds looking at each other on the road endseq

startseq two dogs on pavement moving toward each other endseq

-----Predicted-----

startseq two dogs play with each other in the grass endseq





generate\_caption("1002674143\_1b742ab4b8.jpg")

-----Actual-----

startseq little girl covered in paint sits in front of painted rainbow with her hands in bowl endseq  
startseq little girl is sitting in front of large painted rainbow endseq  
startseq small girl in the grass plays with fingerpaints in front of white canvas with rainbow on it endseq  
startseq there is girl with pigtails sitting in front of rainbow painting endseq  
startseq young girl with pigtails painting outside in the grass endseq

-----Predicted-----

startseq little girl in pink dress is lying on the side of the grass endseq



generate\_caption("101669240\_b2d3e7f17b.jpg")

-----Actual-----

startseq man in hat is displaying pictures next to skier in blue hat endseq  
startseq man skis past another man displaying paintings in the snow endseq  
startseq person wearing skis looking at framed pictures set up in the snow endseq  
startseq skier looks at framed pictures in the snow next to trees endseq  
startseq man on skis looking at artwork for sale in the snow endseq

-----Predicted-----

startseq two people are hiking up snowy mountain endseq

